Creator acquisition on a £400 budget: what worked, what didn't, and what I'd build first

Delza Brown · Campaign run 28 May – 29 June 2026

1. Introduction

Stoke is a subscription accountability app for married men. I needed 10 to 20 creator partnerships to test whether user-generated content could move subscriptions. Market rate for that test was roughly £6,000. I had £400.

The bet: TikTok Shop's affiliates have already self-selected into commission-only pay, so they'd take performance-only pay off-platform too.

They didn't, cleanly. I recruited five creators for £80 each against a £400-plus market rate per partnership. None of them posted. Recruitment worked, activation didn't, and the gap between those two outcomes is the finding this memo is about.

2. What I was testing, and how I'd know

Primary metric: creators who publish content. Not signups, not replies. Everything upstream is a proxy, and treating a proxy as the goal is the specific error this campaign is a case study in.

Success, stated before the run, was 10 to 20 creators producing content. What proof would actually have looked like is eagerness: creators signing up and getting stuck in, not signing up and going quiet. That distinction matters because the failure mode I hit produces the same top-of-funnel numbers as success.

Secondary metrics, in the order I trusted them: opportunities (replies flagged as live conversations), reply rate, sends.

3. Operating tenets

These were fixed before the first send and they explain most of the design decisions below.

Deliverability outranks measurement. No tracking pixels, no links in the first email. Both would have improved visibility into the send-to-reply gap and both would have cost deliverability on cold domains. Reaching the inbox was worth more than knowing what happened inside it.

The offer is fixed, the copy is variable. Three subject and body variants, A, B and C. One offer. If the offer had moved too, nothing would have been attributable.

Judge on the metric closest to the outcome, not the one easiest to read. Reply rate is available instantly and means less. This tenet is the reason I scaled the variant I did.

Every constraint becomes a design decision. £400 ruled out paid partnerships, so the offer had to carry itself.

4. What happened

The offer

$15 flat per sign-up, 20% off for their audience, free access for three months, approved within 48 hours. No product to buy, no shipping, no reimbursement wait, no payout hold. Every friction TikTok Shop imposes on an affiliate, removed. I swapped their percentage for a flat bounty deliberately, it prices cleanly against a subscription and it kills the return-window risk entirely.

The funnel

4,290 emails to 1,466 reachable creators over 20 days. 102 replies. 5 signed. 0 posted. Full table in the appendix.

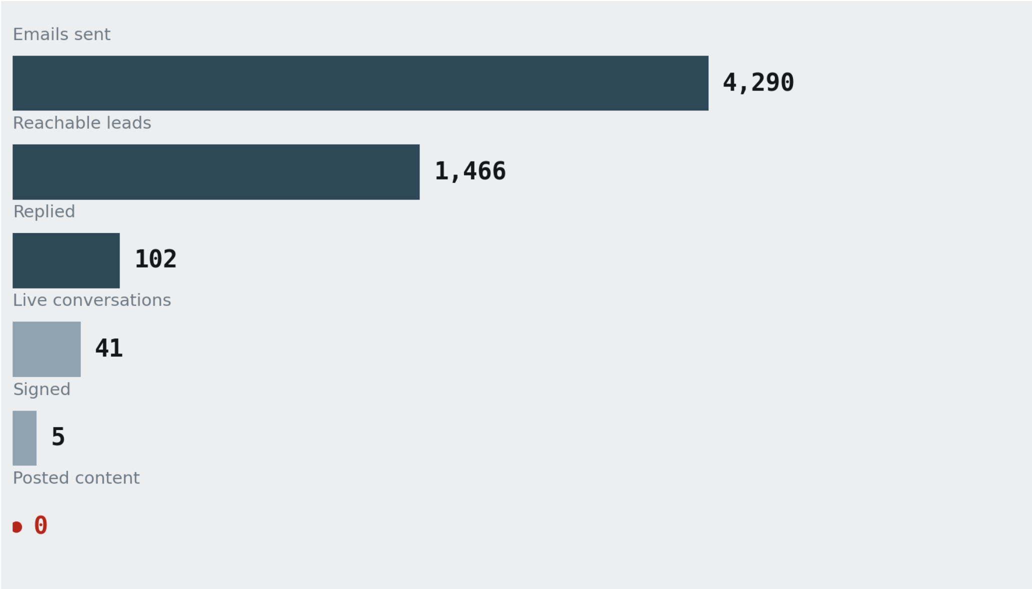

£80 to sign a creator, against a £400-plus market rate for the equivalent outcome the paid route buys. Cost per activated creator isn't high, it's undefined. The denominator is zero.

The scale decision, day 8

Roughly even bases across all three variants in the pre-scale window, about 330 leads each. Raw reply rate has B ahead, 3.38% against A's 2.36%, but that's the weaker signal, it doesn't distinguish genuine interest from a canned rate-card counter dressed as one. Opportunities is the stronger metric, it reflects actual conversational intent, and on that metric, same base, A was already ahead: 10 flagged against B's 8, C's 4.

Those counts are too close to separate statistically on a base of 330. I knew that at the time, which is why I didn't scale on the count alone. I read the flagged replies case by case to check the intent behind them was real. They held up. I scaled A on 4 June, eight days in, as a judgment call the numbers supported rather than a result the numbers proved.

5. Lessons learned

Zero creators posted, and I know why. I built the onboarding, access and a discount code, but I never built what came after it. There was no automated sequence to keep Stoke in front of a creator once they'd signed, no scheduled nudges pushing them from signed up toward actually filming. I was running sourcing, copy, the landing page and the dashboard solo inside a month, and follow-through after signup is where that ran out of runway. It's not a mystery in this data, it's a specific gap, and it's the first thing I'd build if I ran this again.

Raw reply rate would have pointed at the wrong variant. Roughly 60% of replies were canned rate cards. A rate card is a polite no, not a negotiation. Any metric that counts it as engagement will rank the variant that attracts the most polite noes, which is exactly what reply rate did here.

The rails were never the question. The habit was. I knew going in why TikTok Shop affiliates accept commission-only pay. The platform gives them a one-click impulse purchase, nothing between a viewer being sold on a product and owning it, and that converts at a rate which makes a low commission worth their time. Stoke is a subscription, so I couldn't borrow those rails. What I couldn't know without running it was whether the habit travels: whether someone already working commission-only on one platform takes a performance-only offer somewhere else. Underneath that sat a more hopeful bet, that newer creators wouldn't price in the one-click advantage they'd be giving up and would read the two offers as equivalent, while more experienced ones would run the numbers and pass. I can't call that conclusive. Plenty may have said no simply because it wasn't paid, or wasn't their lane.

Supply was a harder ceiling than economics. Every creator on the sourcing platform has a handle. Roughly a fifth have an email attached to it. And the categories that drive TikTok Shop, clothing and cosmetics, pull a creator base that skews female, while Stoke sells to married men. Two cuts, both before spend entered the conversation. That's not a return-on-spend question, it's a supply question, and no amount of budget fixes it.

6. What I'd do next

Build the post-signup sequence first. It's the one gap I can name precisely and the only way to find out whether five signups and silence means weak interest or just nobody following up.

Count the split properly. An exact intrigued-versus-rate-card count per variant. Right now that's a read from going through replies by hand, not a hard number, and the whole metric-choice argument rests on it.

Check the approval gate. Whether any of the 5 signups stalled at the 48-hour review, or whether all five cleared it and still didn't post. Those are different problems with different fixes.

7. What I'm not claiming

No paid-UGC video was ever produced, so there's no counterfactual. I can't say the £6,000 route would have worked better, only that this one didn't get to content.

I'm also not claiming this settles the hypothesis. Proof would have looked like eagerness, and eagerness was the thing I built nothing to sustain, so weak interest and weak follow-up land as the same silence in this data. I can't cleanly separate them.

A1. Funnel

Stage	Count	Rate
Emails sent	4,290	
Reachable leads	1,466	
Replied	102	2.4% of sends
Rate-card counters	~61	~60% of replies
Live conversations	~41	~40% of replies
Signed	5	12% of live
Posted content	**0**	**0%**

A2. Variant performance

	A	B	C
Sent (full run)	2,334	977	979
Replied (full run)	55	33	14
Reply rate (full run)	2.36%	3.38%	1.43%
Leads, pre-scale window	332	334	333
Opportunities, pre-scale	**10**	**8**	**4**

A was scaled on 4 June, which is why its full-run send volume is roughly 2.4 times the others. Reply rates are shown full-run; the scale decision was made on the pre-scale window only.

A3. Cost

Total spend	£400
Market rate for equivalent test	~£6,000
Cost per signup	£80
Cost per activated creator	**undefined (denominator zero)**

A4. Definitions

Reachable lead. A creator with an email address attached to their profile on the sourcing platform, passing the gender filter.

Reply. Any response to any email in the sequence.

Rate-card counter. A reply quoting standard paid rates rather than engaging with the offer. Classified by hand, treated as a decline.

Opportunity. A reply flagged as a live conversation in the sending platform at the time of the run. Counted live, not derived retrospectively.

Signed. Application approved and discount code issued.

A5. Method and limitations

Campaign exports from the sending platform were aggregated in Python, grouped by variant, and reply rates computed from sent and replied counts. Opportunity counts were recorded live during the run rather than reconstructed afterwards.

Open and click tracking were deliberately not implemented in the first email, to protect deliverability. Landing page and app-download link clicks were tracked, but nobody posted, so downstream attribution never became relevant.

The rate-card versus live classification is a manual read, not a rule-based one, and no per-variant count exists yet. Opportunity counts of 10, 8 and 4 on bases of roughly 330 are not statistically distinguishable from one another.

A6. Live artifacts

Source data and pipeline (GitHub): github.com/delzacbrown-ship-it/Stoke-Report